

Research Report: Agents Memory

Generated on February 21, 2026

Agents Memory: An Integrated Survey of Architectural, Computational, and Physical Perspectives

Introduction

The notion of memory lies at the heart of any autonomous or intelligent agent, yet its definition and implementation differ markedly across disciplines. In computer science literature, agent architecture is described as a blueprint that organizes facts, goals, and optionally a plan library into a coherent system (Wikipedia, Agent Architecture). This structural view implicitly treats memory as the repository of facts and goals that guide behavior. Parallel developments in cognitive science have produced architectures such as Soar, which model general intelligence through symbolic working memory and long-term knowledge stores (Wikipedia, Soar). Meanwhile, hardware-level investigations of emerging non-volatile memories, notably HfO_x-based resistive RAM, reveal that variability and bit-error-rate directly affect the reliability of the physical substrate on which agents store information (ArXiv, In-Line Test of Variability). Together, these strands suggest that a comprehensive understanding of agents' memory must span abstract architectural models, algorithmic learning mechanisms, and the underlying physical media. This report synthesizes findings from agent architecture, memory-model research, reinforcement-learning exploration, cognitive-agent theory, and autonomous-system control to delineate the current state of knowledge, practical applications, and open challenges concerning agents' memory.

Main Findings – Architectural Foundations

Agent architectures provide the first layer of memory organization. Contemporary designs treat an agent as a collection of declarative facts coupled with a goal hierarchy, sometimes augmented by a plan library that encodes procedural knowledge (Wikipedia, Agent Architecture). Cognitive architectures extend this view by embedding memory within a unified computational substrate that supports reasoning, learning, and perception. Soar, for example, implements a long-term procedural memory that stores production rules and a working memory that holds the current problem state, enabling the system to achieve general intelligence through iterative problem solving (Wikipedia, Soar). Recent AI-agent surveys highlight a shift toward hybrid architectures that combine large language models (LLMs) with dedicated perception, planning, and tool-use modules, thereby distributing memory across both statistical embeddings and explicit symbolic stores (ArXiv, AI Agents). This evolution underscores a tension between end-to-end learned representations and modular, interpretable memory components, a tension that is reflected in the ongoing debate over the precise definition of "agent"

itself.

Main Findings – Computational Memory in Learning

Reinforcement learning (RL) treats memory primarily as the agent’s internal estimate of value functions and policies, which are updated through interaction with the environment. Traditional optimistic exploration methods provide value bonuses only after a state has been visited, limiting first-visit discovery (ArXiv, Value Bonuses). The Value Bonuses with Ensemble errors (VBE) algorithm addresses this limitation by maintaining an ensemble of random action-value functions whose estimation errors generate first-visit optimism, thereby enhancing deep exploration and reducing sample complexity (ArXiv, Value Bonuses). Complementary work on hyperparameter optimization for RL, exemplified by the ARLBench benchmark, demonstrates that efficient evaluation of learning configurations can be achieved by selecting representative task subsets, implicitly treating the hyperparameter space as a memory of past performance that guides future searches (ArXiv, ARLBench). These studies reveal that memory in RL is both a dynamic statistical construct—updated online—and a meta-level repository that records algorithmic experience for reuse across tasks.

Main Findings – Memory Management in Transactional and Distributed Settings

Beyond learning, agents must manage memory safely in concurrent and mobile environments. In software transactional memory (STM), opacity requires that even aborted (“doomed”) transactions never observe inconsistent states, a guarantee that hinges on robust memory reclamation. Epoch-based reclamation, while preventing access violations, incurs high memory consumption and fails to reclaim non-transactional allocations (ArXiv, Opacity of Memory Management). A novel approach that blends incremental validation with sandboxing mitigates these drawbacks, offering a more parsimonious memory footprint while preserving opacity. In the domain of agent mobility, the FIPA-based interoperable mobility proposal defines a flexible protocol stack that enables agents to migrate across heterogeneous platforms such as JADE and AgentScape (ArXiv, FIPA-based Mobility). This architecture treats the agent’s state—including its memory of facts, goals, and learned models—as a portable payload, emphasizing the need for standardized serialization and deserialization mechanisms to preserve memory integrity during migration.

Main Findings – Physical Memory Substrates

The reliability of an agent’s memory ultimately depends on the characteristics of the underlying hardware. Studies of HfO_x-based resistive RAM demonstrate that spatial and temporal variability across manufacturing lots, wafers, and chips can be quantified through bit-error-rate (BER) as a function of design margin (ArXiv, In-Line-Test). By deriving BER from in-line test data, designers obtain a holistic metric that informs both technology evaluation and product-level error-correction strategies. Importantly, the same BER derivation methodology can be applied to built-in-self-test (BIST) schemes, suggesting that agents could dynamically assess the health of their memory hardware and adapt their behavior—e.g., by invoking redundancy or re-training—when error rates exceed tolerable thresholds. This hardware-aware perspective bridges the gap between abstract memory models and the physical constraints that shape real-world agent performance.

Applications – Cognitive and Autonomous Agents

The convergence of architectural, algorithmic, and hardware insights enables sophisticated applications across cognitive and autonomous domains. In cognitive agents, autonomously active neural networks that maintain self-sustained dynamics while remaining receptive to sensory input have been shown to perform online non-linear independent component analysis, effectively constructing internal memory representations that map sensory components onto attractor relics (ArXiv, Autonomous Neural Networks). When coupled with emotion-driven diffusive control mechanisms, such agents can reduce the dimensionality of policy search spaces, thereby achieving more efficient lifelong utility maximization (ArXiv, Cognition and Emotion). In autonomous vehicle systems, memory manifests in the form of learned models for perception, planning, and control; the control stack relies on accurate state estimation and predictive models stored in memory to execute longitudinal and lateral maneuvers safely (ArXiv, Control Strategies). Moreover, autonomous mobility-on-demand (AMoD) platforms depend on centralized memory of demand patterns, traffic congestion, and fleet status to orchestrate vehicle dispatch, illustrating how large-scale memory aggregation supports real-time operational decision making (ArXiv, AMoD). Across these examples, the ability to store, retrieve, and update information reliably is a prerequisite for intelligent, adaptive behavior.

Challenges – Definition, Integration, and Reliability

A persistent challenge is the lack of a universally accepted definition of “agent,” which hampers the development of standardized memory interfaces (Wikipedia, Agent Architecture). This definitional ambiguity leads to divergent memory models: symbolic fact-goal stores, statistical embeddings in LLMs, and low-level hardware bits, each with distinct semantics and performance characteristics. Integrating these heterogeneous memories raises contradictions; for instance, the high memory consumption of epoch-based reclamation in STM conflicts with the lightweight, portable state required for interoperable mobility. Similarly, the variability of resistive memories threatens the deterministic operation of safety-critical autonomous systems, yet current RL exploration methods assume reliable value updates. Addressing these contradictions demands cross-layer co-design, wherein hardware error models inform algorithmic robustness, and architectural standards prescribe serialization formats that preserve semantic consistency across platforms.

Challenges – Evaluation and Benchmarking

Evaluating agents’ memory capabilities remains problematic. Existing RL benchmarks often neglect memory-intensive tasks such as long-horizon planning or continual learning, focusing instead on short-term reward maximization. The ARLBench initiative mitigates this by providing a curated set of hyperparameter-optimization tasks, yet it still does not capture the full spectrum of memory-related challenges, such as catastrophic forgetting or memory-constrained inference. Cognitive-agent research similarly suffers from a paucity of metrics that jointly assess symbolic reasoning, emotional modulation, and autonomous neural dynamics. Consequently, there is a gap in holistic evaluation frameworks that can compare agents across domains, a gap highlighted by recent calls for balanced metrics encompassing effectiveness, efficiency, robustness, and safety (ArXiv, AI Agents). Developing such benchmarks will be essential for quantifying progress in memory-centric agent design.

Conclusion

Agents' memory emerges as a multi-faceted construct that spans high-level architectural organization, algorithmic learning dynamics, concurrent memory management, and the physical reliability of storage media. Architectural surveys reveal a consensus that agents must maintain declarative facts, goals, and procedural plans, while modern AI systems augment these with learned embeddings and modular perception-planning pipelines. Reinforcement-learning research contributes advanced exploration mechanisms that treat memory as a statistical estimator, and hyperparameter-benchmarking efforts democratize access to memory-aware optimization. Concurrently, advances in transactional memory management and interoperable mobility protocols address safety and portability concerns, whereas hardware studies on resistive RAM provide quantitative tools for assessing memory health. Despite these advances, the field lacks a unified definition of agent memory, suffers from fragmented evaluation practices, and faces reliability challenges when deploying agents in safety-critical autonomous settings. Future research should pursue cross-disciplinary co-design, integrating hardware error models with robust learning algorithms, and establishing comprehensive benchmarks that capture the full lifecycle of memory—from acquisition and consolidation to retrieval and adaptation. Such an integrated approach will be pivotal for realizing agents that can store, reason over, and act upon information with the fidelity required for real-world autonomy.